

Solving Bomberman

WITH Q-TABLE AND NEURAL NETWORKS

BEHROOZ MONTAZERAN

JANNIS HEISING

TOMÁŠ SLÁMA

20. 12. 2023




Introduction

PROBLEM STATEMENT

PROBLEM STATEMENT

- **6 actions** – 4 movement directions + wait + place bomb
- destructible crates, indestructible walls
- **+1** for coin, **+5** for killing opponent



METHODS USED

METHODS USED

Q-learning – optimizing action-value Q-function:

- **Q-table:** store each input/output in a table
- **Neural Network:** learn function via a neural network

METHODS USED

Q-learning – optimizing action-value Q-function:

- **Q-table:** store each input/output in a table
- **Neural Network:** learn function via a neural network

TD update – Use best possible predicted value in next step:

$$Y_t \leftarrow R_t + \gamma \cdot \max_a Q(s_{t+1}, a),$$
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \mu \cdot (Y_t - Q(s_t, a_t)).$$

METHODS USED

Q-learning – optimizing action-value Q-function:

- **Q-table:** store each input/output in a table
- **Neural Network:** learn function via a neural network

TD update – Use best possible predicted value in next step:

$$Y_t \leftarrow R_t + \gamma \cdot \max_a Q(s_{t+1}, a),$$
$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \mu \cdot (Y_t - Q(s_t, a_t)).$$

Stability can be improved with policy/target model:

- **target:** trained on the results of actions
- **policy:** updated slowly / replaced periodically

Evaluation with `elo.py`:

- comparing relative strengths of agents using the elo system
- agents start with 1 000 elo, updated after win/loss/draw

Evaluation with `elo.py`:

- comparing relative strengths of agents using the elo system
- agents start with 1 000 elo, updated after win/loss/draw

Training with `train.py`:

- subcommands for training traits (coin picking, agent hunting)
- periodically calculates strength with `elo.py`
- useful training flags: `--infinite <n>`, `--continue`

TRAINING COMMAND EXAMPLE

TRAINING COMMAND EXAMPLE

```
TASKS = {  
    ...  
    "complete": [  
        # first learn to pick up coins  
        (["--scenario", "coin-heaven", "--n-rounds", "100"], False),  
        # then learn to hunt/evade the rule-based agent  
        (["rule_based_agent", "--scenario", "empty", "--n-rounds", "1000"], False),  
        # then learn to work with crates  
        (["--scenario", "classic", "--n-rounds", "1000"], False),  
        # then learn to work with crates + hard agent  
        (["rule_based_agent", "--scenario", "classic", "--n-rounds", "200"], True),  
        # then learn to work with crates + really hard agent  
        (["binary_agent_v3", "--scenario", "classic", "--n-rounds", "200"], True),  
        # finally play against itself (None gets substituted)  
        ([None, "--scenario", "classic", "--n-rounds", "200"], True),  
    ],  
    ...  
}
```

Neural Networks

Network structure: fully connected layers with ReLU

- uses TD for Q-value updating

Network structure: fully connected layers with ReLU

- uses TD for Q-value updating

Main idea – carefully crafted binary state vector:

1... 5: direction to the nearest **coin**

6...10: direction to the nearest **crate**

11...15: direction to where a bomb will **kill opponent**

16...20: direction to **safety** (if in danger of dying)

21: **can place a bomb** and there is a way to escape it

Network structure: fully connected layers with ReLU

- uses TD for Q-value updating

Main idea – carefully crafted binary state vector:


- 1... 5: direction to the nearest **coin**
- 6...10: direction to the nearest **crate**
- 11...15: direction to where a bomb will **kill opponent**
- 16...20: direction to **safety** (if in danger of dying)
- 21: **can place a bomb** and there is a way to escape it

Calculated by **searching state space** (other players stand still)

- ⇒ finds path even in ongoing explosions
- ⇒ shortest on-board path \neq shortest path stepwise

Q-VECTOR EXAMPLE (1)

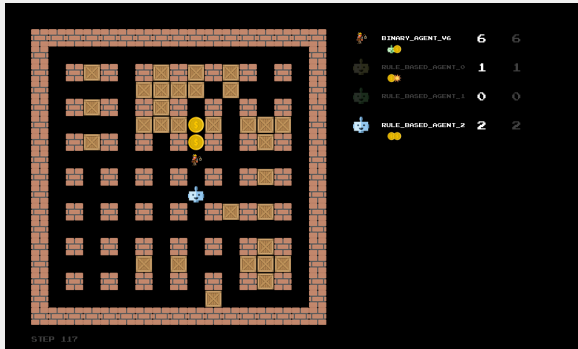
Q-VECTOR EXAMPLE (1)



0	1	0	0	0	0	1	0	0	0	0	0	0	0	0	0	1	0	0	0	0
↓					↓					↓					↓					↓
Nearest coin to the right					Nearest crate to the right (by state space search)					No players reachable					Safety to the right					Can't place a bomb

Q-VECTOR EXAMPLE (2)

Q-VECTOR EXAMPLE (2)



1 0 0 0 0 | 1 0 0 0 0 | 0 0 0 0 1 | 0 0 0 0 0 | 1

↓ ↓ ↓ ↓ ↓

Nearest coin up Nearest crate up Place bomb to endanger player Safe Can place a bomb

TRAINING PARAMETERS

TRAINING PARAMETERS

Hyperparameters

```
BATCH_SIZE = 256
MEMORY_SIZE = 1000

# chances of picking random action
# happens per-round, not per-epoch!
EPS_START = 0.99
EPS_END = 0.05
EPS_DECAY = 10

# (soft) update rate of target network
TAU = 1e-3

# LR of the optimizer
LR = 1e-4

# discount for future states
GAMMA = 0.99

OPTIMIZER = optim.Adam
LAYER_SIZES = [1024, 1024]
```

Event Rewards

```
# hunt coins
MOVED_TOWARD_COIN: 50
DID_NOT_MOVE_TOWARD_COIN: -100

# hunt people
MOVED_TOWARD_PLAYER: 25
DID_NOT_MOVE_TOWARD_PLAYER: -10

# blow up crates
MOVED_TOWARD_CRATE: 1

# basic stuff
e.INVALID_ACTION: -100
DID_NOT_MOVE_TOWARD_SAFETY: -500

# be active!
USELESS_WAIT: -100

# meaningful bombs
PLACED_USEFUL_BOMB: 60
PLACED_SUPER_USEFUL_BOMB: 150
DID_NOT_PLACE_USEFUL_BOMB: -500
```



TRAINING LOOP

TRAINING LOOP

1. `train.py` with the "complete" task and `--infinite` 2
2. pick the best-performing model based on calculated elo
3. decrease `EPS_START`, `EPS_END`, `TAU` and `LR` hyperparameters
4. go to step 1 (with `--continue` to not overwrite the model)

TRAINING LOOP

1. `train.py` with the "complete" task and `--infinite`
2. pick the best-performing model based on calculated elo
3. decrease `EPS_START`, `EPS_END`, `TAU` and `LR` hyperparameters
4. go to step 1 (with `--continue` to not overwrite the model)

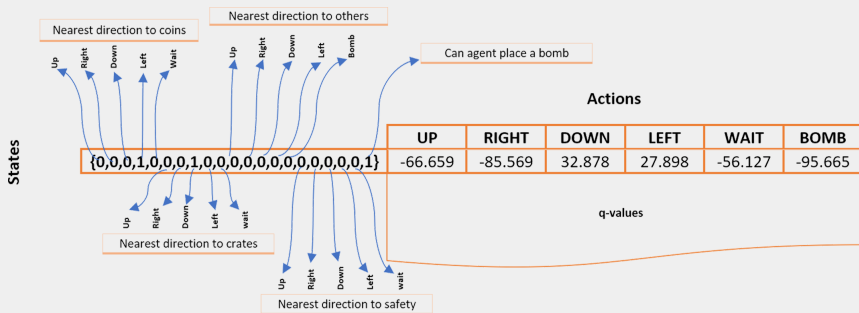


Q-table

STORING Q-TABLE

STORING Q-TABLE

- Theoretically $207 \times 2^{21} \times 6 = 2.60466278e9$ possible states
- 207 free tiles for the agent to go
- 2^{21} possible actions-state pairs
- 6 actions



STORING Q-TABLE

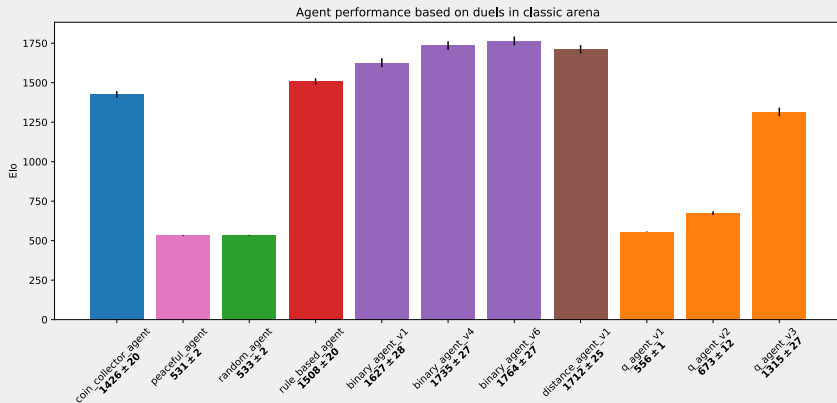
In practice, agent encounters much fewer states

Scenario	Opponent(s)	Number of Rounds	Maximum new encountered states	Encountered states	Run time	Total run time
Coin-heaven	-	5,000	~100	100	~00:03:40	00:03:40
Sparse-crate	-	30,000	~560	660	~00:20:30	00:24:10
Sparse-crate	Peaceful-agent	40,000	~2290	2950	~00:47:35	01:11:45
classic	Peaceful-agent	40,000	0	2950	~00:27:28	01:39:13
classic	Peaceful-agent & rule-based-agent	40,000	~2530	5480	~01:15:17	02:54:30
classic	rule-based-agent & rule-based-agent & rule-based-agent	100,000	~5315	10795	~05:44:10	08:38:40
Doubled	Above in order	255,000	~16	10811	~04:20:40	12:59:20

Conclusion

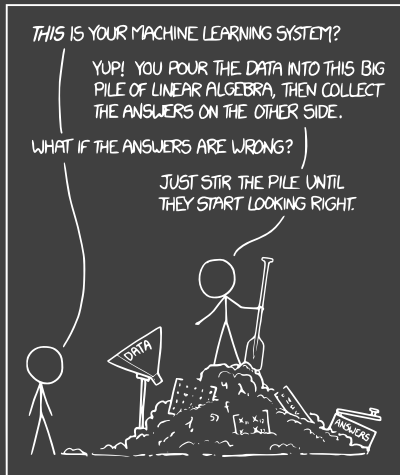
RESULTS

RESULTS



THANKS FOR YOUR ATTENTION!

SPECIAL THANKS TO ULLRICH KÖTHE AND TUTORS



REFERENCES



BEHROOZ MONTAZERAN, JANNIS HEISING, T. S.
BombermanML.

<https://github.com/xiaoxiae/BombermanML>.

Accessed: 2023-11-28.



MUNROE, R.

XKCD 1838.

<https://xkcd.com/1838/>.

Accessed: 2023-11-28.



RE-LOGIC.

Terraria.

<https://www.terraria.org/>.

Played: 2023-12-20.